

Introduction to Web Science/595: Assignment #8

Dr. Nelson

Francis Pruter

Thursday, 13 November, 2014

Contents

Problem 1	3
Problem 2	7
Problem 3	8
Problem 4	9
Problem 5	10
Problem 6	11
Problem 7	12
Problem 8	12
Problem 9	13
Problem 10	16

Problem 1

The goal of this project is to use the basic recommendation principles we have learned for user-collected data. You will modify the code given to you which performs movie recommendations from the MovieLense data sets.

The MovieLense data sets were collected by the GroupLens Research Project at the University of Minnesota during the seven-month period from September 19th, 1997 through April 22nd, 1998. It is available for download from <http://www.grouplens.org/node/73>

There are three files which we will use:

1. `u.data`: 100,000 ratings by 943 users on 1,682 movies. Each user has rated at least 20 movies. Users and items are numbered consecutively from 1. The data is randomly ordered. This is a tab separated list of

```
user id | item id | rating | timestamp
```

The time stamps are unix seconds since 1/1/1970 UTC.

Example:

```
196 242 3 881250949
186 302 3 891717742
22 377 1 878887116
244 51 2 880606923
166 346 1 886397596
298 474 4 884182806
115 265 2 881171488
```

2. `u.item`: Information about the 1,682 movies. This is a tab separated list of

```
movie id | movie title | release date | video release date | IMDb URL | unknown | Action |
```

The last 19 fields are the genres, a 1 indicates the movie is of that genre, a 0 indicates it is not; movies can be in several genres at once. The movie ids are the ones used in the `u.data` data set.

Example:

```
161|Top Gun (1986)|01-Jan-1986||http://us.imdb.com/M/title-exact?Top%20Gun%20(1986)|0|1|0|
162|On Golden Pond (1981)|01-Jan-1981||http://us.imdb.com/M/title-exact?On%20Golden%20Pond
163|Return of the Pink Panther, The (1974)|01-Jan-1974||http://us.imdb.com/M/title-exact?Re
```

3. `u.user`: Demographic information about the users. This is a tab

separated list of:

```
user id | age | gender | occupation | zip code
```

The user ids are the ones used in the u.data data set.

Example:

```
1|24|M|technician|85711
2|53|F|other|94043
3|23|M|writer|32067
4|24|M|technician|43537
5|33|F|other|15213
```

The code for reading from the u.data and u.item files and creating recommendations is described in the book Programming Collective Intelligence (check email for more details). You are to modify recommendations.py to answer the following questions. Each question your program answers correctly will award you 10 points. You must have the question answered completely correct; partial credit will only be awarded if your answer is very close to the correct one.

1. What 5 movies have the highest average ratings? Show the movies and their ratings sorted by their average ratings.

DISCUSSION After downloading recommendation.py from [1], I modified the loadMovieLens function to include the loading of user data (age and gender).

Listing 1: loadMovieLens

```

180 def loadMovieLens():
    # Get movie titles

    for line in open('u.item'):
        (id, title) = line.split('|')[0:2]
185     movies[id] = title

    # Load data
    prefs = {}
    for line in open('u.data'):
        (user, movieid, rating, ts) = line.split('\t')
190     prefs.setdefault(user, {})
        prefs[user][movies[movieid]] = float(rating)

    # Load user data
    userData = {}
    for line in open('u.user'):
195     (user, age, gender) = line.split('|')[0:3]
        userData.setdefault(user, {})
        userData[user]["Ratings"] = prefs[user]
        userData[user]["Age"] = age
        userData[user]["Gender"] = gender
200     return userData

```

Additionally, I loaded all the lists required to answer all the questions (the ones I was able to answer).

Listing 2: main

```

if __name__ == '__main__':
    # get a list of all ratings
    prefs = loadMovieLens()
205 movieRatings = {}
    movieRatingsWomen = {}
    movieRatingsWomenGT40 = {}
    movieRatingsWomenLT40 = {}
    movieRatingsMen = {}
210 movieRatingsMenGT40 = {}
    movieRatingsMenLT40 = {}

    for usr, usrData in prefs.iteritems():
215         for movie, movieRating in usrData["Ratings"].iteritems():
            movieRatings.setdefault(movie, []).append(int(movieRating))
            if (usrData["Gender"] == 'F'):
                movieRatingsWomen.setdefault(movie, []).append(int(movieRating))
                if (int(usrData["Age"]) < 40):
220                     movieRatingsWomenLT40.setdefault(movie, []).append(int(movieRating))
                elif (int(usrData["Age"]) > 40):
                    movieRatingsWomenGT40.setdefault(movie, []).append(int(movieRating))

            else:
225                 movieRatingsMen.setdefault(movie, []).append(int(movieRating))
                if (int(usrData["Age"]) < 40):
                    movieRatingsMenLT40.setdefault(movie, []).append(int(movieRating))
                elif (int(usrData["Age"]) > 40):
230                     movieRatingsMenGT40.setdefault(movie, []).append(int(movieRating))

```

SOLUTION

When computing the highest 5 average movie ratings, there was actually a 10 way tie. The following is the code used to calculate the highest average:

Listing 3: loadMovieLens

```

avgRatings = {}

240 for movie, ratings in movieRatings.iteritems():
    avgRatings[movie] = np.mean(ratings)

avgRatingsDescending = OrderedDict(sorted(avgRatings.items(), key = lambda k: k
    [1], reverse = True))

ctr = 0
245 prev = 0
with open("q1", "w+") as f:
    for x in avgRatingsDescending.items():
        if (ctr < 5 or prev == x[1]):
            f.write(str(x[1]) + ' & ' + str(x[0]) + '\\\\n\\n')
250         prev = x[1]

```

```
ctr = ctr + 1
else:
    break
```

Below is the solution:

Avg Rating	Movie Title
5.0	They Made Me a Criminal (1939)
5.0	Santa with Muscles (1996)
5.0	Someone Else's America (1995)
5.0	Saint of Fort Washington, The (1993)
5.0	Entertaining Angels: The Dorothy Day Story (1996)
5.0	Marlene Dietrich: Shadow and Light (1996)
5.0	Star Kid (1997)
5.0	Aiqing wansui (1994)
5.0	Prefontaine (1997)
5.0	Great Day in Harlem, A (1994)

Problem 2

2. What 5 movies received the most ratings? Show the movies and the number of ratings sorted by number of ratings.

SOLUTION

The following movies had the most ratings:

Num Ratings	Movie Title
583	Star Wars (1977)
509	Contact (1997)
508	Fargo (1996)
507	Return of the Jedi (1983)
485	Liar Liar (1997)

I used the following code segment to calculate the solution:

Listing 4: loadMovieLens

```
numRatings = {}

260 for movie, ratings in movieRatings.iteritems():
    numRatings[movie] = len(ratings)

numRatingsDescending = OrderedDict(sorted(numRatings.items(), key = lambda k: k
265 [1], reverse = True))

ctr = 0
prev = 0
with open("q2", "w+") as f:
    for x in numRatingsDescending.items():
        if (ctr < 5 or prev == x[1]):
270             f.write(str(x[1]) + " & " + str(x[0]) + "\\n")
            ctr = ctr+1
            prev = x[1]
        else:
            break
```

Problem 3

3. What 5 movies were rated the highest on average by women? Show the movies and their ratings sorted by ratings.

SOLUTION

The following is a list of the highest rated movies by women; again, although the problem only calls for the top 5, there was an 11 way tie.

Avg Ratings	Movie Title
5.0	Stripes (1981)
5.0	Someone Else's America (1995)
5.0	Everest (1998)
5.0	Maya Lin: A Strong Clear Vision (1994)
5.0	Mina Tannenbaum (1994)
5.0	Year of the Horse (1997)
5.0	Faster Pussycat! Kill! Kill! (1965)
5.0	Foreign Correspondent (1940)
5.0	Telling Lies in America (1997)
5.0	Prefontaine (1997)
5.0	Visitors, The (Visiteurs, Les) (1993)

I used the following code segment to calculate the solution:

Listing 5: loadMovieLens

```

avgRatings = {}

280 for movie, ratings in movieRatingsWomen.iteritems():
    avgRatings[movie] = np.mean(ratings)

avgRatingsDescending = OrderedDict(sorted(avgRatings.items(), key = lambda k: k
    [1], reverse = True))

285 ctr = 0
prev = 0
with open("q3", "w+") as f:
    for x in avgRatingsDescending.items():
        290 if (ctr < 5 or prev == x[1]):
            f.write(str(x[1]) + ' & ' + str(x[0]) + '\\\\n\\n')
            prev = x[1]
            ctr = ctr + 1
        else:
            break

```


Problem 4

4. What 5 movies were rated the highest on average by men? Show the movies and their ratings sorted by ratings.

SOLUTION

The following is a list of the highest rated movies by men; again, although the problem only calls for the top 5, there was an 15 way tie.

Avg Ratings	Movie Title
5.0	They Made Me a Criminal (1939)
5.0	Santa with Muscles (1996)
5.0	Letter From Death Row, A (1998)
5.0	Saint of Fort Washington, The (1993)
5.0	Quiet Room, The (1996)
5.0	Entertaining Angels: The Dorothy Day Story (1996)
5.0	Marlene Dietrich: Shadow and Light (1996)
5.0	Star Kid (1997)
5.0	Little City (1998)
5.0	Aiqing wansui (1994)
5.0	Prefontaine (1997)
5.0	Love Serenade (1996)
5.0	Leading Man, The (1996)
5.0	Great Day in Harlem, A (1994)
5.0	Delta of Venus (1994)

I used the following code segment to calculate the solution:

Listing 6: loadMovieLens

```

avgRatings = {}

300  for movie, ratings in movieRatingsMen.iteritems():
    avgRatings[movie] = np.mean(ratings)

avgRatingsDescending = OrderedDict(sorted(avgRatings.items(), key = lambda k: k
    [1], reverse = True))

305  ctr = 0
    prev = 0
    with open("q4", "w+") as f:
        for x in avgRatingsDescending.items():
            if (ctr < 5 or prev == x[1]):
310                f.write(str(x[1]) + ' & ' + str(x[0]) + '\\\\ \n' )
                prev = x[1]
                ctr = ctr + 1
            else:
                break

```

Problem 5

5. What movie received ratings most like Top Gun? Which movie received ratings that were least like Top Gun (negative correlation)?

Problem 6

6. Which 5 raters rated the most films? Show the raters' IDs and the number of films each rated.

SOLUTION

Below is a list of the top 5 raters:

Num Ratings	Rater's ID
736	405
678	655
632	13
538	450
516	276

I used the following code segment to calculate the solution:

Listing 7: loadMovieLens

```
numUsrRatings = {}

for usr, usrData in prefs.iteritems():
    numUsrRatings[usr] = len(usrData["Ratings"])

numRatingsDescending = OrderedDict(sorted(numUsrRatings.items(), key = lambda k: k
[1], reverse = True))

ctr = 0
prev = 0
with open("q6", "w+") as f:
    for x in numRatingsDescending.items():
        if (ctr < 5 or prev == x[1]):
            f.write(str(x[1]) + " & " + str(x[0]) + "\\ \\ \\ \n")
            ctr = ctr+1
            prev = x[1]
        else:
            break
```

Problem 7

7. Which 5 raters most agreed with each other? Show the raters' IDs and Pearson's r , sorted by r .

Problem 8

8. Which 5 raters most disagreed with each other (negative correlation)? Show the raters' IDs and Pearson's r , sorted by r .

Problem 9

9. What movie was rated highest on average by men over 40? By men under 40?

SOLUTION

The following movies all received a perfect 5.0 for men over 40:

Avg Ratings	Movie Title
5.0	Indian Summer (1996)
5.0	Leading Man, The (1996)
5.0	Unstrung Heroes (1995)
5.0	Little Princess, The (1939)
5.0	Great Day in Harlem, A (1994)
5.0	They Made Me a Criminal (1939)
5.0	Spice World (1997)
5.0	Boxing Helena (1993)
5.0	Little City (1998)
5.0	Double Happiness (1994)
5.0	Poison Ivy II (1995)
5.0	Two or Three Things I Know About Her (1966)
5.0	Star Kid (1997)
5.0	Ace Ventura: When Nature Calls (1995)
5.0	Grateful Dead (1995)
5.0	Aparajito (1956)
5.0	World of Apu, The (Apur Sansar) (1959)
5.0	Rendezvous in Paris (Rendez-vous de Paris, Les) (1995)
5.0	Prefontaine (1997)
5.0	Solo (1996)
5.0	Late Bloomers (1996)
5.0	Strawberry and Chocolate (Fresa y chocolate) (1993)
5.0	Marlene Dietrich: Shadow and Light (1996)
5.0	Faithful (1996)
5.0	Hearts and Minds (1996)

The following is a list of perfect 5.0 for men under 40

Avg Ratings	Movie Title
5.0	Letter From Death Row, A (1998)
5.0	Perfect Candidate, A (1996)
5.0	Saint of Fort Washington, The (1993)
5.0	Quiet Room, The (1996)
5.0	Magic Hour, The (1998)
5.0	Entertaining Angels: The Dorothy Day Story (1996)
5.0	Maya Lin: A Strong Clear Vision (1994)
5.0	Angel Baby (1995)
5.0	Star Kid (1997)
5.0	Love in the Afternoon (1957)
5.0	Aiqing wansui (1994)
5.0	Prefontaine (1997)
5.0	Love Serenade (1996)
5.0	Leading Man, The (1996)
5.0	Crossfire (1947)
5.0	Santa with Muscles (1996)
5.0	Delta of Venus (1994)

I used the following code segment to calculate the solution:

Listing 8: loadMovieLens

```

avgRatingsGT40 = {}
avgRatingsLT40 = {}

350 for movie, ratings in movieRatingsMenGT40.iteritems():
    avgRatingsGT40[movie] = np.mean(ratings)

avgRatingsDescendingGT40 = OrderedDict(sorted(avgRatingsGT40.items(), key = lambda
    k: k[1], reverse = True))

355 for movie, ratings in movieRatingsMenLT40.iteritems():
    avgRatingsLT40[movie] = np.mean(ratings)

avgRatingsDescendingLT40 = OrderedDict(sorted(avgRatingsLT40.items(), key = lambda
    k: k[1], reverse = True))

360

with open("q9", "w+") as f:
    prev = -1
    f.write("Older than 40:\n")
365 for x in avgRatingsDescendingGT40.items():
    if (prev == -1 or prev == x[1]):
        f.write(str(x[1]) + " & " + str(x[0]) + "\\ \\ \\ \n")
        prev = x[1]
    else:
370         break

    prev = -1
    f.write("Under 40:\n")
375 for x in avgRatingsDescendingLT40.items():
    if (prev == -1 or prev == x[1]):

```

```
        f.write(str(x[1]) + " & " + str(x[0]) + "\\ \\ \\n")
        prev = x[1]
    else:
        break
```

Problem 10

10. What movie was rated highest on average by women over 40? By women under 40?

SOLUTION

The following is a list of perfect 5.0 for women over 40:

Avg Ratings	Movie Title
5.0	Safe (1995)
5.0	Pocahontas (1995)
5.0	Bride of Frankenstein (1935)
5.0	Grand Day Out, A (1992)
5.0	Nightmare Before Christmas, The (1993)
5.0	Gold Diggers: The Secret of Bear Mountain (1995)
5.0	Mary Shelley's Frankenstein (1994)
5.0	Mina Tannenbaum (1994)
5.0	Letter From Death Row, A (1998)
5.0	Band Wagon, The (1953)
5.0	Angel Baby (1995)
5.0	Balto (1995)
5.0	Quest, The (1996)
5.0	Top Hat (1935)
5.0	Wrong Trousers, The (1993)
5.0	Tombstone (1993)
5.0	Foreign Correspondent (1940)
5.0	Best Men (1997)
5.0	Swept from the Sea (1997)
5.0	Ma vie en rose (My Life in Pink) (1997)
5.0	Funny Face (1957)
5.0	In the Bleak Midwinter (1995)
5.0	Shall We Dance? (1937)
5.0	Visitors, The (Visiteurs, Les) (1993)
5.0	Great Dictator, The (1940)
5.0	Shallow Grave (1994)

The following is a list of perfect 5.0 for women under 40:

Avg Ratings	Movie Title
5.0	Stripes (1981)
5.0	Don't Be a Menace to South Central While Drinking Your Juice in the Hood (1996)
5.0	Someone Else's America (1995)
5.0	Grace of My Heart (1996)
5.0	Horseman on the Roof, The (Hussard sur le toit, Le) (1995)
5.0	Wedding Gift, The (1994)
5.0	Heaven's Prisoners (1996)
5.0	Everest (1998)
5.0	Umbrellas of Cherbourg, The (Parapluies de Cherbourg, Les) (1964)
5.0	Nico Icon (1995)
5.0	Maya Lin: A Strong Clear Vision (1994)
5.0	Mina Tannenbaum (1994)
5.0	Year of the Horse (1997)
5.0	Faster Pussycat! Kill! Kill! (1965)
5.0	Telling Lies in America (1997)
5.0	Prefontaine (1997)
5.0	Backbeat (1993)

I used the following code segment to calculate the solution:

Listing 9: loadMovieLens

```

avgRatingsGT40 = {}
avgRatingsLT40 = {}

385
for movie, ratings in movieRatingsWomenGT40.iteritems():
    avgRatingsGT40[movie] = np.mean(ratings)

avgRatingsDescendingGT40 = OrderedDict(sorted(avgRatingsGT40.items(), key = lambda
    k: k[1], reverse = True))

390
for movie, ratings in movieRatingsWomenLT40.iteritems():
    avgRatingsLT40[movie] = np.mean(ratings)

avgRatingsDescendingLT40 = OrderedDict(sorted(avgRatingsLT40.items(), key = lambda
    k: k[1], reverse = True))

395

with open("q10", "w+") as f:
    prev = -1
    f.write("Older than 40:\n")
    for x in avgRatingsDescendingGT40.items():
        if (prev == -1 or prev == x[1]):
            f.write(str(x[1]) + " & " + str(x[0]) + "\\ \\ \\ \n")
            prev = x[1]
        else:
            break

    prev = -1
    f.write("Under 40:\n")
410
    for x in avgRatingsDescendingLT40.items():

```

```
if (prev == -1 or prev == x[1]):  
    f.write(str(x[1]) + " & " + str(x[0]) + "\\ \\ \\ \\n")  
    prev = x[1]  
else:  
    break
```

415

References

- [1] K. Arthur Endsley. recommendation.py. <https://github.com/arthur-e/Programming-Collective-Intelligence/blob/master/chapter2/recommendations.py>, 2012.
- [2] SciPy. Tentative NumPy Tutorial. http://wiki.scipy.org/Tentative_NumPy_Tutorial, 2012.